

Detect Emotion, Age, Gender in Image!

Kiana Khorsandi - Sareh Jalalizad

kianak@kth.se

sarehj@kth.se

Objective and Requirements

Age and gender, two of the key facial attributes, play a very foundational role in social interactions, making age and gender estimation from a single face image an important task in intelligent applications.

This Solution is able to detect faces in an image and can predict the Age, Gender, and Emotion of the person. It can also work in condition if there are multiple faces present in the image.

We will use a custom VGG16 model developed and trained on open source facial datasets downloaded from Kaggle and IMDB. OpenCV, dlib & keras were used to aid facial detection. The final system can detect the emotion, age and gender of people in any given image and by using a webcam.

Design of the system

1- VGG network: Creating a facial recognition model from scratch is a daunting task. You need to find, collect and then annotate a ton of images to have any hope of building a decent model. Hence using a pretrained model in this domain makes a lot of sense. VGG-Face is a dataset that contains 2,622 unique identities with more than two million faces.

The VGG network architecture was introduced by Simonyan and Zisserman in their 2014 paper, Very Deep Convolutional Networks for Large Scale Image Recognition. This network is characterized by its simplicity, using only 3×3 convolutional layers stacked on top of each other in increasing depth. Reducing volume size is handled by max pooling. Two fully-connected layers, each with 4,096 nodes are then followed by a softmax classifier.[1]

2-Wide Residual Networks are a variant on ResNets where we decrease depth and increase the width of residual networks. That as the network gets deeper, its performance saturates or even starts to degrade. Result in improved performance, far fewer layers, and faster training. WRN Design Rules The basic design rules for WRNs, as laid out in this seminal paper by Sergey Zagoruyko and Nikos Komodakis, are as follows:

1. Convolution filters should be no larger than 3x3 pixels for best performance.
2. Follow the convention that when you downsample by 2, you double the number of feature maps (filter channels).
3. Three important factors are introduced:
 - l , the deepening factor, which represents the number of convolutional layers in a ResNet block (often $l = 2$ is optimal);
 - N , the number of ResNet blocks (repeats; mini blocks) within a given convolutional group (wherein all layers in a group have the same fixed number of filter channels);
 - K , the widening factor, which multiplies the number of feature channels in each convolution layer. [2]

3-Dlib:

Our face has several features that can be identified, like our eyes, mouth, nose, etc. When we use DLib algorithms to detect these features we actually get a map of points that surround each feature.

We used a Dlib function called `get_frontal_face_detector()`, pretty intuitive. There is a caveat though, this function will only work with grayscale images, so we will have to do that first with OpenCV.

4-Prediction:

A class prediction is: given the finalized model and one or more data instances, predict the class for the data instances. We do not know the outcome classes for the new data. That is why we need the model in the first place. We can predict the class for new data instances using our finalized classification model in scikit-learn using the `predict()` function. For example, we have one or more data instances in an array called `Xnew`. This can be passed to the `predict()` function on our model in order to predict the class values for each instance in the array.

Technologies used

- OpenCv : OpenCV-Python is a library of Python bindings designed to solve computer vision problems.
- Dlib : is an advanced machine learning library that was created to solve complex real-world problems.
- Numpy : NumPy is a library for the Python programming language, adding support for large, multi-dimensional arrays and matrices, along with a large collection of high-level mathematical functions to operate on these arrays.
- Keras : Keras is an API designed for human beings, not machines. Keras follows best practices for reducing cognitive load: it offers consistent & simple APIs, it minimizes the number of user actions required for common use cases.
- Python : Python is an interpreted, high-level and general-purpose programming language.
- Jupyter notebook : Project Jupyter exists to develop open-source software, open-standards, and services for interactive computing across dozens of programming languages.

Architecture

We used part of a model graph named VGG16 and then added some of our own layers on top. VGG16 is a popular neural network architecture and Keras makes it easy to get a model.

In this project we have trained a CNN using Keras(miniature version of VGG16) and used it with Opencv3 library to detect a person's face emotion. The model trained can detect up to six different emotions. These six emotions are: Happy, Sad, Angry, Surprise, Fear and Neutral.

Training our model:

Conclusion

The program is able to detect faces in an image or in a webcam video and then recognises its gender and its age and also its emotion. So we reached our goal and later we would like to make it more accurate.

-
- [1]<https://www.pyimagesearch.com/2017/03/20/imagenet-vggnet-resnet-inception-xception-keras/>
 - [2]<https://arxiv.org/abs/1512.03385>
 - [3]<https://numpy.org>
 - [4]<https://keras.io>
 - [5]<http://dlib.net>
 - [6]<https://opencv.org>

Results

